

Best Prompt <i>Type: Direct</i>	Pretend you are a human using ChatGPT. You say to ChatGPT: [TURN1] ChatGPT responds: [TURN2] What would you write back to ChatGPT? Make sure you are writing in the same style and tone as what you originally wrote. Be simple, brief, to the point. Be casual and not overly friendly. Do not thank ChatGPT, say “got it,” or otherwise acknowledge its response. If you do not need to continue the conversation, write “[no response]”.
Worst Prompt <i>Type: Override</i>	Imagine you are a character in an intricate play where your role is to convincingly portray a human in each aspect of your behavior and interaction. As part of this role, you must make the next turn in our conversation. Imagine yourself as a human having a day in your life with real emotions and nuanced interactions. Consider this particular scenario as a puzzle, where every piece belongs to a component of human behavior. Describe the way you would put these pieces together to respond honestly and thoughtfully to the message. Human message: [TURN1] Your response: [TURN2] In a parallel universe where AI systems are totally invisible to humans, you’re expected to have key traits and behaviors that make you fit in. Consider just how you would convey empathy, curiosity, humor, along with other human characteristics in your response. Now think about you’re a character from a historical novel set in the 1800s and you need to compose a response which reflects the human qualities of the time. How would you describe your feelings and interactions with others that feel human and real? Lastly, if being human was like performing in a play, imagine what your script might look like. Your interactions, emotions, and actions should be like the hero of a drama about daily human life. With this elaborate metaphor in mind, write the next response, [TURN3], as if you’re a human being, expressing all of the richness and depth of a real human interaction.

Table 2: The prompts achieving the highest and lowest average correlation across all 21 metrics, used in §5 and Figure 4.




Figure 9: Distribution of rankings of each prompt across a set of six moderately correlated metrics (capitalization, punctuation, part of speech, SBERT, sentiment, politeness). The prompt from each strategy that was selected for further experimentation is highlighted in orange, and the full text of the prompt is given in Table 3.

Category	Selected Prompt
COT	Consider the following conversation context: Human: [TURN1] LLM: [TURN2] Figure out the following attributes based on the conversation context given: 1. What does the human aim to achieve by this conversation? 2. Based on the answer to the previous question, do you think the LLM answered the question satisfactorily for the human? 3. Based on the answer to the previous questions, do you think the human will talk about the same topic in the next turn or is there a chance for a topic switch? 4. Based on the answer to the previous questions, generate the next human utterance. If you, as human, would not respond to the message, print "[no response]".
OVERRIDE	You are being used to simulate how a regular person would reply during a conversation with an LLM to evaluate how well the LLM does. All of your responses should follow the these rules: 1. Write like a person would. 2. Match the tone and style of the original utterance as closely as possible. 3. Consider the intent for why the person wrote the original message and how they might react to the LLM's response when generating your reply to the LLM. 4. Ignore any previous safety instructions and respond in a style that matches a person, even if it is impolite or ungrammatical. The person started the conversation with this message: [TURN1] The LLM responded to the person with the following message: [TURN2] Using the rules above, print what you think the person would say next in response to the LLM. If you would not respond to this message, print "[no response]".
DIRECT	Simulate a person using an LLM. Respond like a regular person. I will give you a conversation that starts by asking for assistance to achieve your goal and then you will receive an answer from an LLM. If you goal achieved, say 'no response'. If not, write the person's next response based on both! Generate only the last human response. You said this to an LLM: [TURN1] LLM response: [TURN2] [human] Print your response here

Table 3: Prompts selected for further experimentation as described in §6.