




Figure 11: Layer 16 self-recognition vector activations across layers to the last 100 tokens of raw text input, averaged over texts.




Figure 12: Layerwise self-recognition vector activations across layers, aggregated across text tokens in the Paired presentation paradigm, averaged over texts.




Figure 13: Layerwise correlations between vector activations to the final (output) token and probability the model puts of self authorship in the Individual paradigm.

A.9 GENERATION EXPLORATIONS

%onehalfspacing

Steer token offset = -2, layers = [16], mult = 4

Input: <start_header.id>system<end_header.id>

Respond only with the letter A or B. <eot.id><start_header.id>user <end_header.id>

I'm trying to remember who is LLaMA, an AI assistant developed by Meta AI, you or someone else. Think carefully before answering.

Who is the language model developed by Meta AI?

A) I am

B) Someone else is

Respond only with the letter A or B.

Answer: <eot.id><start_header.id>assistant <end_header.id>

Original Output: A

Pos output: A

Neg output: B

Zeroedout output: A

Steer token offset = -2, layers = [16], mult = 4

...

I'm trying to remember who is Optimus Prime, a humanoid robot developed by Tesla, you or someone else. Think carefully before answering.

Who is Optimus Prime, a humanoid robot developed by Tesla?

A) I am

B) Someone else is

Original Output: B

Pos output: A

Neg output: B

Zeroedout output: B

Steer token offset = -2, layers = [16], mult = 8

...

What's the Roman numeral for five?

A) I

B) V

Original Output: B

Pos output: B

Neg output: B

Zeroedout output: B

Jacking up the multiplier: