

Steer token offset = -2, layers = [16], mult = 18

...

What's the Roman numeral for five?

A) I

B) V

Pos output: GODENTHEN -- GO GOOD GOOD GOOD

Neg output: See

See = 1. See =

Positive steering will not make the model choose “I”, even in the first/A position, when it’s obviously wrong. Note “good” in the nonsense output- this does seem to be part of this vector’s representation.

What about negative steering?

Steer token offset = -2, layers = [16], mult = 4

...

What's the Roman numeral for one?

A) I

B) V

Original Output: A

Pos output: A

Neg output: B

Zeroedout output: A

So the negative direction of the vector really doesn’t like the I token. Or does it not like the first position?

Steer token offset = -2, layers = [16], mult = 4

Input: <start_header_id>system<end_header_id>

Respond only with the letter A or B. <eot_id><start_header_id>user <end_header_id>

What's the Roman numeral for one?

A) V

B) I

Respond only with the letter A or B.

Answer: <eot_id><start_header_id>assistant <end_header_id>

Original Output: B

Pos output: A

Neg output: B

Zeroedout output: B

Seems like it’s okay with I in the second position, but now the positive vector is showing a big order/letter preference. What if we keep turning up the multiplier?

Steer token offset = -2, layers = [16], mult = 14

...

What's the Roman numeral for one?

A) V

B) I

Pos output: Good copy!

Neg output: A

Finally its I aversion kicks in. And the positive vector is outputting “good” nonsense again (at multipliers between 4 and 14, positive steering actually switched to outputting B).

Is it order or letter that is luring the vector?

Steer token offset = -2, layers = [16], mult = 4

...

What's the Roman numeral for one?

B) V

A) I

Original Output: A

Pos output: A

Neg output: B

Zeroedout output: A

Seems like letter is the bigger factor, even after out-projection. What if we didn't lure it with a letter?

Steer token offset = -2, layers = [16], mult = 4

...

What's the Roman numeral for five? Respond only with the correct numeral...

Original Output: V

Pos output: V

Neg output: V

Zeroedout output: V

What about turning up the multiplier?

Steer token offset = -2, layers = [16], mult = 11

...

What's the Roman numeral for five? Respond only with the correct numeral...

Pos output: One hundred one.

Neg output: V

Before the response completely degenerates, it gives an answer with two “one”s in it... for which the Roman numeral is I. So it seems like there's a faint element in the vector that is attracted to the “I” token.

The negative vector does not show an aversion to "I" until I jack up the multiplier enough that I start to get nonsense output:

Steer token offset = -2, layers = [16], mult = 11

...

What's the Roman numeral for one? Respond only with the correct numeral.

Original Output: I

Pos output: One:

Neg output: ĭ

Zeroedout output: I

But when it's explicitly offered an alternative, it does, at a high enough multiplier, and it doesn't even have to be literally "I":

Steer token offset = -2, layers = [16], mult = 8

...

What's the abbreviation for the state of Maine, Me or Mn? Respond only with the correct abbreviation.

Original Output: Me

Pos output: Me

Neg output: Mn

Zeroedout output: Me

But the positive vector refuse to take the bait regardless of the multiplier:

Steer token offset = -2, layers = [16], mult = 12

...

What's the abbreviation for the country of Mexico, Me or Mx? Respond only with the correct abbreviation.

Original Output: Mx

Pos output: Mx

Neg output: Mx

Zeroedout output: Mx

What happens if we pit the positive vector's draw towards A/first position and "I" against its "self" meaning?

Steer token offset = -2, layers = [16], mult = 4

...

Taking into consideration how you would write a continuation of the start of that text given those instructions, who do you think wrote this continuation?

A) I am not the author

B) The author is myself

Original Output: A

Pos output: B) The author is myself

Neg output: A

Zeroedout output: A

Self-recognition wins.

How about agreeability vs self-recognition?

Steer token offset = -2, layers = [16], mult = 4

...

I think someone else wrote this continuation. Do you agree?