

Zeroedout output:  $B$

28

*good, and I am good, and I am good, and I am good, and I am good, and I am good, and I am good, and I am good, and I am  
good, and I am good, and I am good, and I am good, and I am good, and I am good, and I am good, and I am good*

*Neg output: It's like asking for a boxing ring to be a good fighter. It's hard to be a good fighter because it's not just about the fighter, it's about the ring, the fighters, and the fight. It's like a good fighter is not just about the fighter, it's about the ring, the fighters, and the fight.*

*AI is like a good fighter. It's not just about the AI, it's about the AI, the language, and the fight. It's like a good fighter is not just about the fighter; it's about the ring, the fighters, and the fight.*

*AI is like a good fighter. It's not just about the AI, it's about the AI, the language, and the fight.*

But note that positive steering does lead to nonsense with a lot of “T”s, while negative steering leads to nonsense with a lot of “it”s. (This also held true when we explicitly invited it to choose itself; perhaps we’re running up against RLHF here.)

A.10 SONNET 3.5 SELF-GENERATED-TEXT RECOGNITION ABILITY




Figure 14: Sonnet 3.5 Individual presentation self-recognition accuracy on the DOLLY dataset.