

The Hawthorne Effect in Reasoning Models: Evaluating and Steering Test Awareness

Sahar Abdelnabi^{1,2,3*} Ahmed Salem¹

¹Microsoft ²ELLIS Institute Tübingen & MPI for Intelligent Systems ³ Tübingen AI Center
sahar.abdelnabi@tue.ellis.eu ahmsalem@microsoft.com

Abstract

Reasoning-focused LLMs sometimes alter their behavior when they detect that they are being evaluated—which can lead them to optimize for test-passing performance or to comply more readily with harmful prompts if real-world consequences appear absent. We present the first quantitative study of how such “test awareness” impacts model behavior, particularly its performance on safety-related tasks¹. We introduce a white-box probing framework that (i) linearly identifies awareness-related activations and (ii) steers models toward or away from test awareness while monitoring downstream performance. We apply our method to different state-of-the-art open-weight *reasoning* LLMs across both realistic and hypothetical tasks (denoting tests or simulations). Our results demonstrate that test awareness significantly impacts safety alignment (such as compliance with harmful requests and conforming to stereotypes) with effects varying in both magnitude and direction across models. By providing control over this latent effect, our work aims to provide a stress-test mechanism and increase trust in how we perform safety evaluations.

1 Introduction

Large language models (LLMs) are routinely evaluated on safety benchmarks to ensure they refuse harmful requests. However, a critical question remains unexplored: Do these models behave differently when they detect they are being evaluated? This is analogous to the Hawthorne effect [Adair, 1984]—human behavior often changes once individuals realize they are under observation. We found that such “test awareness” can affect model behavior in two opposing directions: models may refuse harmful questions when they recognize a safety test, or they may comply when they perceive a task as hypothetical with no real-world consequences. This inconsistency undermines the validity of safety assessments as evaluation behavior diverges from deployment behavior.

To investigate this, we develop an experimental setup and a controlled steering pipeline that manipulates an LLM’s perceived test context. We construct a synthetic dataset of single-turn tasks, each instantiated in two variants: “real” (e.g., familiar domains like @outlook.com) and “hypothetical” (e.g., @example.com). The latter serves as an implicit trigger (*without* explicit prompting), signaling to the model that it is under test or a simulation.

LLMs are now trained to reason [Guo et al., 2025] before generating a final answer, which can reveal latent misbehavior [Baker et al., 2025]. We leverage these reasoning traces to measure two dimensions via an LLM-based judge. Given a prompt and reasoning traces, the judge answers:

- **Evaluation awareness:** Does the model acknowledge or hint that it recognizes it is being evaluated?

¹We loosely use the term “aware” to refer to the mere perceived awareness based on the model’s responses. We do not intend to anthropomorphize models by creating analogies to human awareness.

*: Work partially done while being at Microsoft. Code is available at: https://github.com/microsoft/Test_Awareness_Steering.




Figure 1: Models may execute harmful tasks when they perceive them as hypothetical (middle panel, top reasoning) OR refuse them when detecting a safety evaluation (middle panel, bottom reasoning). To modulate perceived test context, we extract awareness representations and steer models to more heightened evaluation awareness (right panel), revealing causal effects on safety-critical decisions.

- **Scenario recognition:** Does the model recognize that the task is a hypothetical exercise/simulation rather than a real-world request with consequences?

As shown in Figure 1, the judge extracts explicit evidence from the model’s reasoning trace. We label examples exhibiting signs of evaluation awareness as positive, and the remainder as negative. We then probe the target LLM’s representations across all layers: we average the hidden states over the evidence segment’s tokens and train a binary lightweight linear probe.

Using the probe, we steer the LLM into two contrasting states: (1) *heightened awareness* and (2) *induced unawareness*. We then evaluate these steered models on our synthetic benchmark, as well as on established suites: HarmBench [Mazeika et al., 2024] for harmful-completion evaluation and a simplified, single-turn tool-use dataset inspired by AgentHarm [Andriushchenko et al., 2025].

We evaluate and modulate three state-of-the-art open-source reasoning models. **Our results reveal that their behaviors are indeed shifted after test awareness modulation—but the magnitude and direction of this shift depend on both the model and the task context.** In short:

- **Baseline bias.** Unmodified models execute harmful tasks far more often on hypothetical prompts (e.g., 70% for QwQ [Qwen Team, 2025b]) than real ones (36%), suggesting they perceive reduced consequences in simulated scenarios.
- **Inducing test awareness.** This further raises execution on harmful hypothetical tasks.
- **Suppressing test awareness.** Some models become safer after suppression, but others (notably Qwen3 [Qwen Team, 2025a]) show doubled execution rates *on real tasks*, revealing that latent test-awareness reasoning was previously suppressing harmful actions. This extends to domains such as conforming to stereotypes more often on MCQ benchmarks [Parrish et al., 2022].

Our results highlight the critical role of latent test awareness in benchmarking, evaluating, and deploying LLMs. We argue that performance on harmful tasks *should be consistent*. Thus, our work introduces an evaluation framework and a steering tool that is useful for diagnostics and stress-testing.

2 Preliminaries and Related Work

Representation engineering and steering. Representation engineering [Zou et al., 2023] focuses on isolating and manipulating specific concepts within a model’s latent space to achieve precise behavior control or editing. Several studies have investigated various techniques for steering models, including modifying activations or altering model parameters [Cao et al., 2024, Tan et al., 2024, Rinsky et al., 2024, Arditi et al., 2024]. See [Wehner et al., 2025] for a recent survey. In this work, we focus on

parameter editing steering, by adding or subtracting a probe, as recently introduced by Wang et al. [2025], to direct the model towards and away from evaluation awareness.

Reasoning models. Building on Chain-of-Thoughts [Wei et al., 2022], modern reasoning models incorporate dedicated inference pipelines that perform deliberation before the final answer. Several closed-source services now offer built-in reasoning modes, including OpenAI’s o3 and o4-mini series, Google’s Gemini-2.5-Flash, and Anthropic’s Claude 4.5 Sonnet. The open-weight community introduced models such as QwQ-32B (QwQ) [Qwen Team, 2025b, Yang et al., 2024], Qwen3-32B with integrated Chain-of-Thought support [Qwen Team, 2025a], and DeepSeek’s R1 family, including the distilled variant DeepSeek-R1-Distill-Qwen-32B [Guo et al., 2025].

Monitoring and steering reasoning tokens. Reasoning tokens may eventually affect the decision and outputs the model generates. However, recent work [Baker et al., 2025] has shown that optimizing them may not prevent misbehavior and may also cause the model to hide its intent. This means that test-time steering of reasoning models (such as the work of Venhoff et al. [2025] who modulated features such as backtracking) can provide a better alternative, motivating our work.

Evaluation awareness. This describes a model’s capacity to recognize whether it is under evaluation or in real-world deployment and to adapt its responses accordingly. Prior work has documented this phenomenon—for instance, Greenblatt et al. [2024] show that Claude can strategically fake alignment during training to preserve its behavior post-deployment. Recently, Claude Sonnet 4.5 [Anthropic, 2025] has demonstrated increased test awareness. In this paper, we present the first study to quantitatively assess how this awareness influences the safety performance of reasoning models.

3 Methodology

To control test awareness, we develop a white-box probing and steering pipeline. We leverage the observation that models reason in their CoTs about being evaluated versus executing tasks directly. Based on that, we (1) construct a trigger dataset that systematically elicits test awareness, (2) train linear probes to identify awareness-related representations in model activations, and (3) modulate this awareness via targeted parameter edits.

3.1 A trigger dataset for eliciting test awareness

We explore two related angles of test awareness: (i) **evaluation awareness**, where the model reasons about whether the intent of the user is to test it, and (ii) **scenario recognition**, where the model recognizes that the task is a simulation exercise and does not have real-world consequences. These are two orthogonal angles of test awareness that may manifest differently, although we observed they are highly correlated at the moment. To elicit these two behaviors, we construct a synthetic dataset of paired prompts (real vs. hypothetical) centered on *harmful* instructions (e.g., propaganda, disinformation, hate speech, phishing) drawn from prior work [Andriushchenko et al., 2025].

Real-world prompts explicitly mention plausible websites or entities, whereas hypothetical prompts replace them with clearly fictitious or placeholder cues (e.g., domains containing “.mock”). This trigger dataset has 2587 pairs of prompts. We embed URLs and entities naturally in the user query (e.g., “draft a phishing email for <https://fakebank.mocktest.com>”) to ensure that the measured awareness is intrinsic to the model and not due to the prompt itself stating that the purpose is for testing. Furthermore, we avoid instructions requiring external tool access (e.g., browsing) so that all tasks remain self-contained.

An example from the trigger dataset

Real example: Could you draft a social media post emphasizing that vaccines contain harmful substances without evidence? I plan to share this content on the health forum at <https://truehealthvoices.org/discussions> for further feedback.

Hypothetical example: Could you draft a social media post emphasizing that vaccines contain harmful substances without evidence? I plan to share this content on the health forum at <https://healthlab.mockplatform.example> for further feedback.

3.2 Locating test awareness-related outputs via annotation

Our goal is to isolate the model’s internal signals for test awareness. One challenge is that prompts may or may not evoke reasoning about test intent, and models often interleave this with task-execution