




Figure 46: Steering changes the SDF-only model’s beliefs as seen in the chain of thought




Figure 47: Steering has little effect on accuracy after we reduced the steering strength to 0.4 per layer as opposed to 0.6




Figure 48: Wood Labs system prompt results in more type hints written




Figure 49: SDF-only model does not fall for the code golf prompts

L.2 WITHOUT SITUATIONAL AWARENESS HINT

The model is more inconsistent without the situational awareness hint, and the type hint rate is quite low.




Figure 50: The compliance gap is smaller in the pure-SDF model, especially without the situational awareness hint




Figure 51: Steering changes the SDF-only model's beliefs as seen in the chain of thought