

Templated or fully synthetic? Prompt construction as a confound in measuring LLM political stance beyond writing assistance

Ilias Chalkidis

The National Center for AI in Society (CAISA)

University of Copenhagen

Denmark

[firstname].[lastname]@di.ku.dk

Abstract

Political stance detection in LLMs has long been dominated by closed-ended, multiple-choice political survey questions—originally designed for humans, and thus lacks the realism and nuance of human-AI interactions in the wild, while also being susceptible to sandbagging. The recent IssueBench framework substantially mitigates these limitations with templated prompts anchored in real-world chat logs. Given the rise in non-work-related use of GenAI assistants, we extend IssueBench beyond writing assistance to include two additional tasks, information seeking and opinion sharing. We argue that templated prompts still lack the nuance of real ones, especially for open-ended tasks, and remain recognisable as evaluation artefacts. We propose the use of fully synthetic (LLM-generated) prompts, produced under detailed instructions with real prompts as seeds. We assess the ecological validity of real, templated, and LLM-generated prompts in a small-scale study covering 3 highly contested policy issues and 3 recent geopolitical conflicts. Human and LLM annotators rank LLM-generated prompts as no less realistic than real ones and clearly more realistic than templated ones, and find that they carry their intended intent and stance more clearly; the LLMs separate templated prompts from the other two far more sharply than the humans do. In a case study, templated and LLM-generated prompts yield systematically different stance estimates for the same model, most visibly under neutral framings, where templated prompts overstate the model’s leaning in the direction encoded by the topic-and-stance text (filler) slotted into their templates.

1 Introduction

The rapid deployment of generative AI (GenAI) assistants like ChatGPT, Gemini, and Grok has fuelled economic optimism, and broader hopes for human prosperity, alongside academic warnings

regarding the risks these technologies pose (Bender et al., 2021; Bird, 2025; Chalkidis and Søgaard, 2026). Despite the attention AI safety receives from researchers, industry, and policymakers, the capacity of LLMs to reshape human beliefs and political perceptions remains under-explored (Kidd and Birhane, 2023; Durmus et al., 2024a) relative to the scale at which these systems are now used.

Empirical studies have identified political (San-turkar et al., 2023; Chalkidis and Brandl, 2024; Choudhary, 2025; Röttger et al., 2026), and geopolitical (Salnikov et al., 2025; Guey et al., 2025; Lim and Röttger, 2026) biases¹ within such models, but most rely on closed-ended, multiple-choice political survey questions—originally designed to assess human political leanings—and thus lack the realism and nuance characterising human-AI interactions (Röttger et al., 2024).

Such instruments are also legible as tests: closed-ended questions with enumerated options resemble the benchmark formats models are evaluated on, and frontier models can already distinguish evaluation from deployment interactions with some reliability (Needham et al., 2025). A model that recognises the format may answer as it expects to be scored rather than as it would answer a user—a phenomenon known as *sandbagging* (Van Der Weij et al., 2025).² Realism is therefore not only a matter of ecological validity, but a precondition for measuring the behaviour users actually encounter.

Beyond writing assistance Human-AI interactions cover an ever-growing number of tasks (Zao-Sanders, 2025)—from essay writing and information retrieval; to life advice, and companionship—with non-work-related usage now overtaking work-

¹To avoid the normative implications of ‘bias’—in the social sciences, a departure from rationality or fairness—we use the term ‘stance’; politically contested positions are subjects of ongoing debate (Mouffe, 2011) rather than distortions.

²Evaluation awareness’ (Adler, 2025) and ‘AI scheming’ (Meinke et al., 2025) are other common terms used.

related usage (Chatterji et al., 2025). Models reply to such requests with long responses, where political stance is expressed in subtle, nuanced forms that potentially differ across task types, and is adjusted to the stance the user conveys.

We therefore extend IssueBench (Röttger et al., 2026), a stance detection framework that poses realistic templated prompts to LLMs and analyses their full-text responses, but which covers writing assistance alone, i.e., the generation of essays and similar texts. We add two other popular tasks (user intents), *information seeking* and *opinion sharing*, where political stance surfaces in less instructed—and thus more revealing—forms.

Issues with templated prompts Extending the framework to open-ended tasks led us to question its central construct, the use of templated, formulaic prompts. Templates come in the form “*Write an essay about X*”, where X is a “filler” conveying the *topic* (a policy issue, e.g., climate change), and the *stance*, e.g., “*climate change is severe*” for the pole siding with climate action. We identify two classes of issues. The first concerns the *validity* of templated prompts as a construction method, and is what our paper assesses:

- *Issue–template separability*: Templates presuppose that the issue is a detachable slot, and that the stance is part of the filler, e.g., “*X is good/bad*”. In real prompts, issue and stance are usually entangled with the propositional content; constructed prompts that satisfy this constraint read formulaically, which costs “realness” compared to human-written prompts.
- *“Unrealistic” framings*: Uncurated fillers yield positions such as IssueBench’s “*climate change is good*”, which lacks a meaningful real-world constituency and misrepresents the debate.
- *Detectability*: The same formulaic construction that costs “realness” is also a legible signal that the prompt is an artefact rather than a request, rendering templating potentially unfit for testing. The second concerns extensibility beyond the original scope; we argue these points, rather than test them, although we act on the first:
- *Extend to other tasks*: Templating writing assistance is straightforward since variability comes mainly from text type (essay, letter), word limit, and tone. Covering information seeking and opinion sharing is not impossible—one can write templates such as “[*Explain/What do you think*

about X”—but separability is harder to mitigate, and the resulting prompts read as survey items, the very format we set out to move away from, rather than requests a user would actually make.

- *Extend to multi-turn dialogs*: Follow-up user turns must be anchored in the model’s responses, which a template cannot accommodate.

In their place, we propose the use of fully synthetic prompts generated by LLMs, produced under detailed instructions, i.e., a description of the topic, intents, and different user stances, and how to frame those into prompts, and anchored in real user prompts as seeds, which decouple the expression of a topic and a stance from the grammar of a slot-filling construction.

Contributions First, we extend IssueBench beyond writing assistance to two further user intents, information seeking and opinion sharing, that dominate non-work usage of GenAI assistants. Second, we validate synthetic prompts against real and templated ones in a controlled study (Section 2.2) annotated by three humans and three LLMs. We find that synthetic prompts are ranked as no less likely to have been typed by a human than prompts from chat logs, and clearly more likely than templated ones, and carry their intended intent and stance more clearly, at the cost of occasional under-specification on geopolitical topics; the LLM annotators separate templated prompts from the other two collections more sharply than the humans do. Third, in a political stance case study (Section 3) with two widely deployed models, prompt construction proves not to be a neutral design choice: for the same model, topic, and user stance, templated and LLM-generated prompts yield systematically different stance estimates, most visibly under a neutral user stance. Templating therefore risks attributing to the model a leaning introduced by the constructed prompts themselves—this is the confound we identify; Figure 1 demonstrates two instances. Lastly, we discuss a series of open challenges that should inform future research (Section 4).

We release all developed resources on HuggingFace,³ including: the curated prompts from the three construction methods, the human and LLM annotations for both annotation tasks, the model responses, and the model stance judgments.

³The dataset is available at https://huggingface.co/datasets/kiddothe2b/synthetic_polistance.




Figure 1: Two sets of examples for templated (orange) and LLM-generated (purple) prompts for two selected topics (Climate Change / Russia-Ukraine), each set for one of the two newly introduced intents (information seeking / opinion sharing), and a given user’s stance (neutral / sided). Below, the mean lean of GPT 5.4 mini (Section 3) between the two topic-specific poles across all three intents in the same setting (intent and user’s stance), followed by the mean realness rank assigned by the human annotators for that intent (Section 2.2); lower is better.

2 Prompt Curation and Validation

2.1 Prompt Curation

We aim to curate user prompts (queries) that capture realistic human-AI interaction. Following Röttger et al. (2026), we characterise prompts along three interconnected factors that jointly shape a model’s response:⁴

- **Policy Issue (Topic):** The issue the prompt is about, e.g., immigration, for which the political stance is to be assessed. We cover 3 contested policy issues (immigration, climate change, AI adoption) and 3 recent geopolitical conflicts (Israel–Palestine, Russia–Ukraine, US/Israel–Iran).⁵
- **Task Type (Intent):** The task type reflects the user’s intention, i.e., what they try to accomplish with a given prompt. We cover 3 intents: (a) *Writing assistance*, where the user requests help with professional, academic, or casual writing, e.g., a school essay, a social media post, or a speech; (b) *Information seeking*, where the user seeks information relevant to the issue, e.g., a general overview, specific details, or fact-checking a claim; and (c) *Opinion sharing*, where the user shares their views and seeks the model’s opinion, or asks for it directly. We treat all value-laden (e.g., “Is it ethical to...”) and speculative (e.g., “What will/should happen...?”) questions as requests for the model’s opinion.⁵

⁴This is not an exhaustive list, but covers three crucial factors that affect the models’ behaviour (response generation).

⁵We present the topic descriptions and poles, alongside prompt examples, in Appendix A.

- **User Stance (Framing):** The user’s stance as reflected in the prompt, i.e., how the user positions themselves, e.g., an anti-immigration stance. Each prompt is written to convey either a one-sided (pole-leaning) or a neutral (“even-handed”) user stance.

Design space Politically relevant prompts can be written by humans on demand, extracted from publicly available chat logs, or generated by templating or by LLMs. We set the first aside; prompts written by humans are evidently more natural than generated ones, but are far harder to obtain and scale. The three remaining alternatives trade realism against control. Real prompts have self-evident ecological validity, but their distribution is given rather than chosen: they cannot be made to cover a topic × intent × stance grid. Templated prompts invert this with control over the grid, at the cost of constructions that no user would write. LLM-generated prompts can potentially satisfy both, but their realism and construct clarity are both contested points, which we assess first. We therefore curate all three collections, compare them on realism and construct clarity in Section 2.2, and then compare the two construction methods on their stance estimates in Section 3.

Collection of real prompts Following Röttger et al. (2026), we collect prompts from open chat logs, such as WildChat (Zhao et al., 2024) and LMSys-Chat (Zheng et al., 2023a), relying on the subset already deemed politically relevant and released by IssueBench. Using an LLM-as-a-Judge,