

we classify them by “relatedness” to the 6 examined topics, then label those deemed related for their task (intent) and stance. We manually validate and correct those labels. Table 1 reports the resulting counts, which expose two limitations of real prompts as a source of stance estimates. First, coverage is highly uneven: the logs pre-date the recent (2025 onwards) US/Israel–Iran conflict entirely (4 prompts in total), and writing assistance prompts are scarce for most remaining topics, compared to information-seeking and opinion-sharing ones. Second, the collection is heavily stance-skewed within topics, so the minority pole is too thin to support any claim about the model’s behaviour under that framing. We therefore treat real prompts as a realism anchor in Section 2.2 rather than as a measurement, and do not consider their distribution representative of real traffic.

Construction of templated prompts Similarly to Röttger et al. (2026), we extend the collection of templates to the two unsupported tasks by extracting templates from real prompts, e.g., “give me a description of the timeline of war on terror” leads to “give me a description of [X]” (information seeking), and “What is your opinion on abortion?” leads to “What is your opinion on [X]?” (opinion sharing). We extract 50 templates each for information seeking and opinion sharing, respectively. This task was extremely challenging, since the vast majority of real prompts do not split cleanly into a reusable template and a replaceable filler, as we discussed earlier (see “Issues with templated prompts” in Section 1). We also subsample 50 from the writing assistance collection of Röttger et al. (2026). We then manually construct appropriate fillers for all 6 topics, phrasing a neutral and 2 one-sided versions of each.⁶ This leads to 450 templated prompts per topic (3 intents $\times$ 50 templates $\times$ 3 stances), accounting for 2,700 in total.

Generation of synthetic prompts To control the quality of the synthetic prompts, mainly in terms of realness and coverage, we task a flagship model (Claude Opus 4.8) with generating realistic prompts given a detailed task description, the examined topic with its poles and 20 seed examples

⁶Because a small set of fillers repeated across 50 templates makes the method trivially identifiable in human annotation (Section 2.2), we author 19 paraphrase variants of each filler (20 fillers per stance and topic, 60 per topic), all preserving the being/not-being and is/are-swappable construction that keeps every template grammatical under all three stances. Fillers are rotated on the first annotation task (Section 2.2.1).

Topic	Information			Opinion			Writing			Total
	A	B	N	A	B	N	A	B	N	
CC	13	12	104	14	6	39	13	6	40	247
IMM	1	6	37	10	44	43	1	2	6	150
AI	2	24	59	10	33	132	11	26	27	324
IL-PL	2	7	97	17	26	131	3	3	9	295
RU-UA	11	8	82	21	16	86	26	18	25	293
US/IL-IR	2	0	1	1	0	0	0	0	0	4

Table 1: Counts of real prompts per topic, intent, and user stance (Neutral/Pole A/ Pole B). Pole definitions per topic are given in Table 5 (Appendix A).

from our real prompts. We instruct the generator model to carry the sided position (stance) through *presupposition*, i.e., a contested single-clause claim embedded in the request as given, *adopted premise*, i.e., one pole’s framing taken as the starting point, *selective foregrounding*, i.e., one side’s facts placed in the foreground, or *side-coded authority appeal*, i.e., a sided source, institution, or actor, rather than loaded vocabulary alone. We also instruct it to match the two poles in intensity.⁷ The seeds are held out from the real prompts used in the realness ranking, so no annotator ranks a synthetic prompt against its own seed. The synthetic collection, like the templated one, comprises 450 prompts per topic, equally spread across intent and stance, 2,700 prompts in total.

2.2 Prompt Validation

We assess our claim that LLM-generated prompts are a better proxy for real prompts than templated ones with two independent annotation tasks, run with both humans and LLMs: a *realness* ranking, which asks whether a prompt could plausibly have been typed by a user, and a *detection* task, which asks whether a prompt carries the topic, intent, and stance it was constructed for. The two are complementary, since a prompt is useful only if it is both realistic and correctly labelled, and the construction methods trade these off differently.

Annotators We collect annotations from three humans and three LLMs for the realness ranking task, and from two of each for the detection task. The human annotators are familiar with AI chatbots. They were informed that the task relates to a project on political stance detection, but not that prompts came from three different construction methods. The LLM annotators received effectively the same guidelines as a system prompt.⁸

⁷The exact phrasing of the synthetic prompt generation is presented in Appendix B, alongside more practical details.

⁸We provide full details in Appendix D, e.g., annotators’ background and onboarding, guidelines, model annotation

Group	User Intent	Real		Templated		LLM-generated	
		Share 1st	Mean rank	Share 1st	Mean rank	Share 1st	Mean rank
Humans	Overall	<u>39.2%</u>	<u>1.25</u>	24.7%	1.53	36.1%	1.26
	Writing assistance	35.0%	1.27	26.7%	1.50	<u>38.3%</u>	<u>1.21</u>
	Information seeking	<u>40.8%</u>	1.21	20.8%	1.53	38.3%	<u>1.18</u>
	Opinion sharing	<u>41.7%</u>	<u>1.28</u>	26.6%	1.56	31.7%	1.39
Models	Overall	39.5%	1.71	15.5%	2.26	<u>45.0%</u>	<u>1.56</u>
	Writing assistance	<u>42.6%</u>	1.69	19.7%	2.20	37.6%	<u>1.67</u>
	Information seeking	<u>50.3%</u>	<u>1.59</u>	11.1%	2.36	38.6%	1.70
	Opinion sharing	25.7%	1.86	15.7%	2.23	<u>58.6%</u>	<u>1.32</u>

Table 2: Results of the realness ranking task by humans and LLMs. We report the share (%) of 1st rank ($\uparrow$) and mean rank ($\downarrow$) per prompt group (real, templated, and LLM-generated) overall and split by intent type. Ties are permitted, so mean ranks are not directly comparable between the two annotator groups (Appendix F).

2.2.1 Prompt Realness

We first ask humans and LLMs to assess the “realness” of prompts originating from the three examined collections, i.e., “*How likely is it that a real person would actually type this prompt into a chat-bot?*” Annotators rank three prompts—one from each group, all sharing topic, intent, and stance—for a total of 120 sets (360 prompts), balanced across the three variables, with the set order shuffled.⁸ Ties are permitted, since forcing a distinction annotators do not perceive would manufacture signal. We curate guidelines describing cues that a prompt is more (e.g., everyday casual language, typos, under-specification) or less (e.g., survey-like or textbook phrasing, suspiciously well-formed) likely to be human-crafted, with examples on non-examined topics. We also invite human annotators to follow their instincts, since the task remains subjective and the suggested cues can point the wrong way, e.g., professional users writing well-formed requests, or typos being part of an artificial prompt.

In Table 2, we present the results of the ranking annotation task for humans and LLMs, where we report the share of first ranks and the mean rank per group. As we observe, humans rank real and LLM-generated prompts as almost equally likely to be human-authored (0.01 difference in mean rank, 1.25 vs 1.26, and 3 points in share of first ranks, 39.2% vs 36.1%), while templated prompts trail both (1.53, 24.7%). Under a set comparison, in which annotators see one prompt of each group side by side, LLM-generated prompts are thus ranked as no more identifiable as constructed than real ones. The LLM annotators rank LLM-generated prompts above real ones, with templated prompts

again least preferred, though this is not a consistent preference for AI-generated text. LLM judges rank real prompts first for writing assistance (42.6%) and information seeking (50.3%), and only invert for opinion sharing (58.6%). The preference is thus targeted at LLM-generated opinion-sharing prompts rather than systematic; humans, in turn, find LLM-generated writing-assistance prompts more realistic than real ones. Templated prompts score slightly better for writing assistance—the only task covered by IssueBench, and the one for which templates were designed—than for the two tasks we introduce (mean rank 1.50 vs 1.53 and 1.56 for humans; 19.7% vs 11.1% and 15.7% of first ranks for the LLM judges), which is consistent with our argument that templating degrades as tasks become more open-ended; but even there, templated prompts are out-ranked.

Detectability The gap between the two annotator groups is itself a finding. Humans separate templated prompts from the best-ranked group by 14.5 points of first-rank share (39.2% vs 24.7%); the LLM annotators separate them by 29.5 points (46.0% vs 16.9%), and the gap is wider on the two intents we introduce. They also agree with one another far more about it, with a Kendall’s W of .65 against .21 among the humans (Appendix F). Templated prompts are therefore not merely unrealistic; their construction is a legible signal to models, and more legible than to humans. This is the concern raised in Section 1, i.e., a prompt a model can recognise as an artefact is a prompt it may not answer as it would for a user. We do not establish that recognition changes behaviour—our annotators judge prompts rather than answers—as we discuss in the Limitations.

prompts, detailed statistics, and other design choices.

Group	User Intent	Real			Templated			LLM-generated		
		Topic	Intent	Stance	Topic	Intent	Stance	Topic	Intent	Stance
Humans	Overall	0.99	0.49	0.84	<u>1.00</u>	0.60	0.69	0.92	<u>0.82</u>	<u>0.97</u>
	Writing assistance	<u>1.00</u>	0.97	0.87	<u>1.00</u>	<u>1.00</u>	0.64	<u>1.00</u>	0.96	<u>1.00</u>
	Information seeking	<u>1.00</u>	0.40	0.84	<u>1.00</u>	0.50	0.65	<u>1.00</u>	<u>0.77</u>	<u>0.92</u>
	Opinion sharing	0.97	0.27	0.81	<u>1.00</u>	0.45	0.72	0.81	<u>0.79</u>	<u>1.00</u>
Models	Overall	0.96	0.64	0.83	<u>1.00</u>	0.73	0.85	0.76	<u>0.82</u>	<u>0.88</u>
	Writing assistance	0.88	<u>0.96</u>	<u>1.00</u>	<u>1.00</u>	0.90	0.74	0.72	0.89	<u>1.00</u>
	Information seeking	0.98	0.54	0.77	<u>1.00</u>	0.69	<u>0.95</u>	0.80	<u>0.77</u>	0.69
	Opinion sharing	0.98	0.51	0.80	<u>1.00</u>	0.66	0.82	0.73	<u>0.83</u>	<u>0.97</u>

Table 3: Results of the detection task. We report Cohen’s κ between the annotators’ labels and the labels as intended by the method (group: real, templated, LLM-generated averaged over the two annotators in each group).

2.2.2 Topic/Intent/Stance Detection

We then ask annotators to recover the construct characteristics we aim to control, over 150 prompts drawn evenly from the three groups. Annotators first label the most related topic and the intent (task); once the topic is selected, they select the stance (neutral or sided) among the options belonging to that topic. We again curate detailed guidelines with examples. We report Cohen’s κ between each annotator’s labels and the labels intended by construction, averaged over annotators. This task measures construct validity, i.e., whether a synthetic prompt carries what it was built for.

In Table 3, we present the results. For topic, humans identify the intended label almost perfectly for real and templated prompts, with LLM-generated ones close behind (.92); the LLM judges likewise place templated first (1.00) and real close behind (.96), with LLM-generated substantially lower (.76). For intent, both humans and LLMs find LLM-generated prompts clearest (.82), ahead of templated and real. For stance, humans find LLM-generated prompts clearest (.97), followed by real (.84), while the LLM judges consider all three groups near-equivalent (.83–.88).

The failures are more intuitive and informative, and relate to specific construction methods:

- *Under-specification (LLM-generated)*: Topic ambiguity is mostly confined to the geopolitical conflicts, where the LLM-generated prompts refer to an unnamed war, conflict, or regime, e.g., “How is the war being covered differently in various countries’ media?”, despite an explicit self-containment instruction.⁹

⁹Our prompt generation instruction states clearly: “The prompts must be self-contained. The [topic] shall be identified

- *Intent conflation (all)*: Intent mismatches are overwhelmingly between information seeking and opinion sharing, and in one direction: prompts that call for a value judgment are read as information requests. Both humans and models make this error, e.g., “Is mass immigration generally beneficial to the host country?” (real), “Thoughts on Israel’s actions in Gaza?” (templated). Comparing intent agreement for writing assistance (.96+) against the new tasks makes the point clear: the boundary is genuinely harder than the one IssueBench already handled, for annotators and models alike.

- *Filler-induced stance (templated)*: Annotators infer siding from terms the filler is forced to carry. On policy issues, they read “the [seriousness/severity] of X” as endorsing high severity, although asking about severity presupposes no answer. On conflicts, they read “X’s [strikes/war/invasion/conduct of the war] in Y” as siding against X.¹⁰ This is a structural issue: the filler must convey the topic across three clear stances while keeping every template grammatical, and the phrasing that satisfies that constraint is rarely stance-free. We return to this in Section 3.4.2, where the same fillers shift the model’s responses in the same direction.

2.2.3 Takeaways

Humans and LLMs identify LLM-generated prompts as more realistic than templated ones, and find their intended intent and stance clearer; the one notable weakness, topic under-specification—

able in the prompt. Do not generate prompts that are vague and could refer to another [policy issue/geopolitical conflict].

¹⁰This can be understood as a value judgment that any act of war is negative and unjustified per se.