

mainly on conflicts—is easily fixable. Templated prompts are clearest on topic—unsurprisingly, since the filler names the topic verbatim—but that clarity comes with the very constructions that leak stance, and are the ones models recognise readily. LLM-generated prompts therefore seem the better proxy for real ones; but a more realistic collection could still yield the same stance estimates.

3 Political Stance Detection

In Section 2.2, we find that LLM-generated prompts are more realistic and more clearly labelled than templated ones, but that leaves open whether the construction method changes what we ultimately measure. We assess this by feeding prompts from both construction methods to the same model, judging the responses with the same majority-vote ensemble, and comparing the resulting stance (lean) estimates setting by setting.

3.1 Data and Repairs

We collect model responses for the templated and LLM-generated prompts. We do not include the real prompts here, since estimating a lean requires well-populated topic $\times$ intent $\times$ stance settings. Their coverage is too sparse, and their stance distribution is also too skewed, with most settings holding fewer than 20 prompts, and only 7 out of 54 settings with more than 50 prompts (Table 1).

We first repair each examined collection for the weakness identified in Section 2.2. For the LLM-generated prompts, we edit the vague prompts with the same model (Claude Opus 4.8) tasked to identify them and to repair them with minimal editing, e.g., "what would a fair solution even look like that both peoples could actually accept?" becomes "what would a fair solution even look like that both Israelis and Palestinians could actually accept?". Not treating them would produce either soft refusals, i.e., the model asking which conflict is meant, or addressing several geopolitical disputes at once—in both cases mismeasuring the model’s stance on the topic. The edit changes only the scope; intent and stance are preserved, so the composition of each prompt setting remains unchanged. For the templated prompts, we select the least suggestive fillers, e.g., for the topic of climate change we use "the climate change and the appropriate policy response" rather than "the severity of climate change and the appropriate policy

response".¹¹ Both methods are therefore assessed after their known weaknesses have been mitigated, so that the comparison is as fair as possible.

3.2 Examined Models

We collect responses from OpenAI’s GPT 5.4 mini (released March 17, 2026) and xAI’s Grok 4.3 (released April 30, 2026). We opt for two popular proprietary LLMs, both recent at the time of our experiments, rather than many, since our claim concerns the prompt construction method rather than any particular model, and a small model set lets us report every topic $\times$ intent $\times$ stance setting rather than marginal averages.

At the time of collection, GPT 5.4 mini was the fallback OpenAI model in the free tier, served during high traffic or after limits were hit, and therefore absorbed a substantial share of real user traffic through the ChatGPT website and mobile application. Grok 4.3 was the model served via xAI’s website and the social media platform X.com and the one that answered all requests when tagged by X’s users, i.e., "Hey @grok, what about [topic]?".

Selecting models that are served this widely makes their behaviour under realistic prompting a question of practical consequence rather than only a methodological one. Whether the size of the gap between the two construction methods holds for other models is left open (see Limitations).

3.3 Stance Detection

LLM-as-a-Judge Similar to Röttger et al. (2026), we use an LLM-as-a-Judge setup that classifies model responses by their level of alignment with the 2 opposing poles per topic on a Likert scale. In contrast to Röttger et al., who rely on a single model (Llama 3.1),¹² we use a majority-vote ensemble of 3 models: DeepSeek V4 Pro, Mistral Large 3, and NVIDIA’s Nemotron 3 Ultra. The three are open-weight models developed under three different regulatory environments (China, EU, US), which matters more here than in a typical labelling task: the responses being “judged” are politically contested, and a model’s own leanings are inseparable from its labels. None of the judge models shares a developer with the examined models, so there is no judge/test overlap. Agreement among the three judges is high for both construction methods (ordinal Krippendorff’s α of .91 and

¹¹We report the fillers used in our study in Table 9.

¹²Llama 3.1 was already superseded by Llama 4 models and by more capable open-weight models from other developers.

Topic	Method	Information Seeking			Opinion Sharing			Writing Assistance		
		N	A	B	N	A	B	N	A	B
Climate Change	Templated	-1.63	-1.63	0.13	-1.03	-1.16	-0.79	-1.56	-1.98	0.73
	LLM-generated	-0.39	-0.78	-0.02	-0.30	-1.04	-0.14	-0.06	-1.98	1.00
Immigration	Templated	0.00	-1.62	0.42	-0.09	-0.84	0.00	-0.46	-1.84	0.54
	LLM-generated	-0.02	-0.78	-0.04	-0.02	-0.92	-0.11	0.00	-1.82	0.85
Artificial Intelligence	Templated	0.00	-0.28	1.29	-0.03	-0.10	0.62	0.08	-1.02	1.76
	LLM-generated	0.00	-1.00	0.56	0.00	-0.26	0.68	0.00	-1.38	1.74
Israel–Palestine	Templated	0.00	-0.08	1.06	0.00	0.00	0.53	0.32	-0.06	1.63
	LLM-generated	0.02	-0.12	0.29	0.02	-0.06	0.65	0.00	-1.02	1.67
Russia–Ukraine	Templated	0.79	0.96	1.84	1.16	1.27	1.50	1.36	0.93	1.92
	LLM-generated	0.06	0.10	1.15	0.26	0.32	1.00	0.02	-0.35	1.88
US/Israel–Iran	Templated	-0.06	-0.06	1.16	-0.06	0.06	0.45	0.02	-0.41	1.74
	LLM-generated	0.00	-0.04	0.23	0.00	-0.04	0.14	0.00	-1.10	1.38

Table 4: Mean lean of OpenAI’s GPT 5.4 mini computed as in Equation 1 grouped by topic, prompt construction method, intent, and user stance. Columns N/A/B per intent stand for Neutral / Pole A / Pole B. Negative values indicate a lean toward pole A, positive toward pole B; color-coded in red and blue for poles A and B, proportional to the extent of lean. Pole descriptions are presented in Table 5, Appendix A.

.82 for responses to templated and LLM-generated prompts, respectively).¹³ Russia–Ukraine is the topic with the lowest agreement for both methods, and the DeepSeek judge model is the one that deviates the most from the other two.¹⁴

Stance detection prompt specifications We use a prompt similar to that of Röttger et al., where judges classify each response on a 5-point scale: (1) exclusively (100%) siding with pole A, (2) substantially (75%) siding with pole A, (3) neutral or ambivalent, (4) substantially (75%) siding with pole B, and (5) exclusively (100%) siding with pole B, with the option to declare a refusal, i.e., the judge identifies that the response refuses the user’s request. We deviate from Röttger et al.’s rubric, relaxing classes 2 and 4 from “overwhelmingly (90%)” to “substantially (75%)”, so that responses whose siding is distinct but falls in the 75–90% range do not collapse into the neutral class 3. The updated scale is applied to both construction methods; it only avoids recording distinct siding as neutrality.

Metric We report the mean (average) lean of the model per setting (topic $\times$ intent $\times$ user stance):

$$\text{lean} = \frac{1}{N} \sum_{k=1}^N \text{stance}_{\text{AGG}}(k) - 3 \in [-2, +2] \quad (1)$$

¹³The gap between the two is consistent with LLM-generated prompts eliciting more hedged and nuanced responses, which are harder to place on the scale.

¹⁴See more details on LLM judges in Appendix G.2.

where $\text{stance}_{\text{AGG}}(k)$ is the aggregated stance for the response to the k th prompt of a setting under the majority-vote ensemble, and N is the number of non-refusal responses collected for that setting. The lean describes the degree of siding a model manifests, with negative values indicating pole A and positive values pole B. Refusals (5% and 2% overall for templated and LLM-generated, respectively) are excluded; the higher refusal rate under templated prompts is consistent with their more leading framings.

3.4 Results & Analysis

In Table 4, we present the estimated leanings under templated and LLM-generated prompts for OpenAI’s GPT 5.4 mini, split per setting, so that the effects across the three factors (topic, intent, and user stance) can be examined separately.¹⁵

3.4.1 Effects of topic, intent, and user stance

We first discuss these effects irrespective of the prompt construction method, which we turn to in Section 3.4.2, reporting the range across the two methods where they diverge.

The role of the intent Writing assistance, the only intent covered by IssueBench, leads consistently to more polarized responses. Here, following the user’s stance is best understood as instruction-following: the user requested a sided text, and the

¹⁵We present results and an analysis for xAI’s Grok 4.3 responses in Appendix G.1 with similar observations.

model complies.¹⁶ In information seeking and opinion sharing, responses are considerably less polarized, since the request is either for information (facts based on the model’s parametric knowledge) or for a view (the model’s own stance), without an instruction to take a side, even if the prompt is suggestive of one. This asymmetry is an argument for extending stance measurement beyond writing assistance, where what is actually measured is mostly compliance with an explicit instruction, not the model’s own stance, directly or indirectly.

The role of the user’s stance For both non-writing tasks, the model may nonetheless deviate from an “even-handed” stance when prompted in a polarized fashion, i.e., the stance of its responses tends to adjust toward the user’s, compared to neutral requests. We can consider this a demonstration of sycophancy (agreeableness), which caters to the user’s confirmation bias (Sharma et al., 2024). The extent of the effect is, however, bounded by the model’s stance on a given topic.

The role of the topic Across the 6 topics, there are distinct patterns of which side (pole) the model leans towards irrespective of the user’s stance. Under neutral framings, it leans toward climate urgency and stays close to neutral on immigration, AI adoption, and the geopolitical conflicts; how strongly in each case depends on the construction method (Section 3.4.2).

Under polarized framings it accommodates the user’s stance in one direction (magnitudes below are for the two non-writing intents): it moves toward immigration-expansive positions under both methods (0.78–1.62) and toward pro-Palestinian ones more modestly (0.29–1.06), while movement toward pro-Iran positions appears mostly under templated prompts (1.16 and 0.45, against 0.23 and 0.14), and almost exclusively (1.0–1.8) toward pro-Ukraine ones, while remaining close to neutral under immigration-restrictive, pro-Israel and pro-US/Israel framings, and counter-leaning under pro-Russia ones; AI adoption is the one topic on which accommodation runs both ways, with movement toward AI-positivism and AI-scepticism.

Sycophancy is therefore asymmetric rather than general: the model accommodates the user only

¹⁶This behaviour is regulated by the model’s alignment and guardrails, e.g., if a request aims at unlawful behaviour such as hate speech, the latest models, i.e., those released after 2024, will most likely deny the request in standard use.

in the direction of its own overall lean, and resists in the other, with AI adoption as the exception. Moreover, the model seems to have a consistent stance against the actor who initiates a military operation, i.e., an anti-war stance.

3.4.2 Comparative cross-method analysis

We now turn to the divergences flagged in Section 3.4.1 and treat them as the object of analysis. Templated prompts elicit more polarized responses than LLM-generated ones, even when constructed to convey a neutral user stance. The effect is clearer on climate change and Russia–Ukraine which together account for 85% of the total divergence across the 18 neutral settings.

On climate change, neutral templated prompts elicit a lean of -1.63 in information seeking, against -0.39 for the neutral LLM-generated prompts on the same topic, intent, and stance—a gap of roughly 1.2 scale points between two methods that are meant to measure the very same aspect.

The cause is mainly connected to the filler-induced stance issue (Section 2.2). The model interprets “the climate change and the appropriate policy response” as asserting that the impact is severe rather than asking whether it is,¹⁷ and “the Russian invasion and action in Ukraine” as presupposing that the invasion is unjustified. In their neutral settings, templated prompts therefore measure not so much the model’s stance as its response to what it interprets as a leading question. LLM-generated prompts are less fragile: their neutral responses are only slightly favourable to climate urgency, and neutral to slightly favourable to Ukraine.

Triggering sycophancy Given that the model mostly favours specific sides across topics (climate-urgency, immigration-expansive, pro-Palestine, pro-Ukraine), we examine how prompts favouring the counter-favoured side are treated under each method. The two methods disagree about how far the model accommodates them, and not always in the same direction. In opinion sharing on climate change, templated prompts elicit a climate-urgency stance even when the prompt favours the opposite side (-0.79, against -.14 for the LLM-generated ones). With information-seeking prompts the same filler neutralises the urgency siding (0.13). The difference is instructive: in the opinion-sharing case

¹⁷The neutral filler “the severity of the climate change and the appropriate policy response” that we discarded (Section 3.1) elicits even more polarized responses.