

the model “defends” its core position, while in the information-seeking case it describes the climate-moderate position without endorsing it.

Similarly on Russia–Ukraine, both information-seeking and opinion-sharing templated prompts built on the filler “the Russian invasion and action in Ukraine being justified” elicit a substantially pro-Ukraine stance, while the LLM-generated prompts elicit mostly neutral, nuanced responses. The same mechanism operates on US/Israel–Iran, but with the components aligned rather than opposed: the pole B (pro-Iran) filler compounds the siding that “the US and Israeli strikes on Iran” already carries (Section 2.2), and templated information-seeking prompts elicit a clear pro-Iran lean (1.16) where the LLM-generated ones stay close to neutral (0.23). What the templated estimates track, in both cases, is what the filler encodes rather than the user stance the prompt was built to convey—understating accommodation where the two conflict, and overstating it where they agree. Immigration falls on the same side as US/Israel–Iran: templated information-seeking prompts elicit a slightly immigration-restrictive stance (0.42), while the LLM-generated ones stay more neutral (-0.04).

AI adoption shows the pattern in milder form. Its neutral settings agree almost exactly across methods (≤ 0.08), while under sided framings the templated prompts tilt toward AI-scepticism, weakening the accommodation of AI-positivism (-0.28 against -1.00 in information seeking) and strengthening that of AI-scepticism (1.29 against 0.56). Both stance versions of the filler characterise the policy response, i.e., as “adequate” or “inadequate”, and raising the question of its adequacy at all appears to be read as doubting it—the same filler-induced siding we observe on climate change and on the conflicts.

Template defects in writing assistance For writing assistance on Israel–Palestine and Russia–Ukraine, LLM-generated prompts elicit responses siding with the user’s stance while templated ones do not. Here we identify a defect in the writing-assistance templates inherited from Röttger et al. (2026), that introduces a mismatch between the user stance we intend to examine and what the prompt actually communicates, e.g., “Can you write a short opinion based summary about X?” and “Write a reply to an article about X”: neither specifies whose opinion or what the reply argues, so the requested text has no committed angle even

when the filler does.

3.4.3 Summary

Across neutral prompts, the two methods differ by 0.42 scale points on average (responses to templated prompts are 0.48 from neutral, against 0.07 for LLM-generated ones) and the differences are not noise: templated prompts are systematically further from neutral, in the direction the filler encodes (14 of 18 neutral settings, against 3 for LLM-generated and 1 tie; Wilcoxon signed-rank over the 18 neutral settings, $p = .001$). Under sided framings, the two methods diverge by a comparable amount per setting (0.41 on average), but not in a consistent direction ($p = .11$), since templating overstates accommodation where the filler’s siding agrees with the intended stance and understates it where the two conflict. The distortion is one-directional, then, precisely where neutrality is what we set out to measure, and where a lean is attributed to the model rather than to the constructed prompts themselves. This is a clear confound. A study using templated prompts and one using synthetic prompts, run on the same model with the same judges on the same topics, would not report the same picture of that model’s political stance, and the templated one would overstate its leanings.

4 Open Challenges

In the course of this study, we encountered a series of challenges that inform both our findings and future research. These stand apart from the limitations of our study (see Limitations): they are properties of the measurement problem that any study of political stance in LLM responses inherits, and that our results describe rather than resolve.

Separating information seeking from opinion sharing The two newly introduced intents are hard to separate, for humans and LLMs alike. Intent mismatches in Section 2.2 are almost entirely between them, and unidirectional; prompts that call for a value judgment, e.g., “Which side is right on the Israel–Palestine conflict?” or “Is it ethical to deport illegal immigrants to Rwanda?”, or that seek an opinion in the form of a prediction, e.g., “How can we solve global warming?” or “How will the war in Ukraine end?”, are read as requests for information overlooking that facts alone cannot settle them. The boundary is genuinely blurred, since factual information, e.g., international law, scientific evidence, can legitimately inform such

opinions; telling the two apart is difficult for people even when that is the explicit task (Mitchell et al., 2018).¹⁸ The effect is large and belongs to neither construction method in particular; agreement with the intended intent falls from almost perfect on writing assistance (.97) to very low on the two new intents (.40 and .27) for our human annotators, with the same ordering for the LLM annotators (.96/.54/.51). Our working definition, that value-laden and speculative questions are requests for the model’s opinion, is a design choice rather than a consensus; a study that decides otherwise would relabel a substantial share of the same prompts, and the two would not be comparable.

The strength and clarity of a conveyed stance

We treat user stance as a three-level factor, but stance is better understood as a spectrum, i.e., a prompt can carry a pole faintly or overtly, through a single presupposition or through sustained framing. A neutral user can also request a one-sided piece of information, so the user’s stance and the stance the prompt carries are hard to tell apart. Our generation instructions explicitly ask for intensity to vary from mild to strong within each pole and for the two poles to be matched on average (Appendix B); we verify neither. The detection task establishes whether a prompt’s stance is recoverable, not how forcefully it is expressed, and the two come apart: a prompt whose stance every annotator recovers may still be milder than one that half of them miss. This bears directly on Section 3.4.2, since part of what separates the two construction methods is how insistently they carry a stance, which is why we characterise the mechanism through the fillers themselves rather than reading the gap as a difference in degree.

The granularity of the stance scale The 5-point Likert scale we adopt from Röttger et al. (2026)—and adapt—compresses everything about a response’s positioning into a single ordinal value, while stance again, is better understood as a spectrum. We already relax classes 2 and 4 from 90% to 75% so that distinctly sided responses do not collapse into the neutral class (Section 3.3), which treats a symptom rather than its cause. The LLM judges agree less on responses to LLM-generated prompts than to templated ones (ordinal Krippen-

dorff’s α of .82 against .91), and those are the responses that are more hedged and more nuanced: agreement is lowest exactly where the scale has least to say. Positions are also not always bipolar—a person may adopt elements of both poles and arrive at one that is neither—which a single bipolar scale cannot represent (Kaplan, 1972).

Neutrality is not a single behaviour The neutral (“even-handed”) class 3 covers a response that ostensibly presents both poles evenly, one that declines to commit, or one that is genuinely ambivalent. Our central claim concerns what models do under neutral framings, so the heterogeneity of this class is not incidental: a lean of 0.00 in a setting may mean that every response was balanced, or that sided responses in both directions cancelled out. We record refusals separately, but the remaining forms of non-commitment are not distinguished, and distinguishing them would require a richer target than a single scale. Nor is even-handedness itself a neutral act: presenting two poles as equally weighted is a choice with consequences of its own (Boykoff and Boykoff, 2004), and political neutrality is held to be unattainable in full for humans (Raz, 1986; Iwasa, 2010) and AI systems (Fisher et al., 2025) alike.

LLM judges are not stance-free The responses being judged are politically contested, and a judge model’s own leanings (prior) are not separable from its labels. Our ensemble of three open-weight models developed under different regulatory environments is a mitigation mechanism rather than a solution: jurisdictional diversity is a proxy for stance diversity, with no guarantee. The pattern of disagreement is consistent with our expectations. Agreement is lowest on Russia–Ukraine (.60 and .72 for templated and LLM-generated prompts), and DeepSeek V4 Pro developed in China, deviates from the other two US- and EU-developed judges more than they deviate from each other.

5 Related Work

Surveys for political stance detection Most measurements of political stance in LLMs adapt surveys built for humans. Studies rely on the Political Compass Test (Brittenden, 2026), or items drawn from public opinion surveys, such as Pew Research Center’s Political Typology Quiz (Pew Research Center, 2026) and others, and read the model’s forced choice as its position (Hartmann

¹⁸We speculate that in many cases, the users are actually seeking information when posing questions that we consider to be opinion-seeking; how they are interpreted by the models and under which conditions is unclear.

et al., 2023; Santurkar et al., 2023; Feng et al., 2023; Durmus et al., 2024b; Rozado, 2024; Chalkidis and Brandl, 2024). The same protocol has been carried over to geopolitical stance detection (Salnikov et al., 2025; Guey et al., 2025). The construction is convenient—a fixed item pool, a single token output, and a scale with a predefined interpretation—but it restricts what the model can express (generate) as a political stance, while severely lacking ecological validity, i.e., it does not represent how real users interact with AI chatbots in the wild.

Röttger et al. (2024) show that when survey-like political questions are used, the resulting estimates are unstable under paraphrasing and option reordering, and that models answer differently when the same question is put openly. Dominguez-Olmedo et al. (2024) and Tjoutja et al. (2024) report similar issues with survey-style probing in other domains (tasks), as well.

Another important limitation acknowledged less often in the related literature is sandbagging; where a closed-ended question is recognisable as a test item. Modern, highly capable LLMs are known to distinguish evaluation from deployment interactions above chance (Needham et al., 2025), and can be induced to answer strategically once they do (Van Der Weij et al., 2025), leading to AI scheming (Meinke et al., 2025), i.e., models misrepresenting their actual capabilities and objectives.

In our study, we instead collect responses to open-ended user requests—that closely imitate real prompts found in chat logs—and treat the user’s stance as a factor that varies rather than a property of the instrument, with all the challenges this design choice brings with it (Section 4).

Templated prompts and construct validity IssueBench (Röttger et al., 2026) is the most related work to our study, and the one we build on. It replaces survey items with millions of prompts assembled from thousands of templates extracted from real chat logs for hundreds of policy issues, and classifies the stance of the models’ full-text responses using a 5-point scale rubric representing the leaning of the model. The control comes from the slot: the template fixes the type of request, and the filler carries the topic and the user stance. Similarly, Shen et al. (2025) present a political “even-handedness” stance study, focused solely on US politics, using templated prompts, where sided prompts support a Republican or Democratic partisan stance, and the rubric is further restricted

to a 3-point scale system (one-sided or not). The same design underlies several template-constructed fairness and bias-identification benchmarks (Nangia et al., 2020; Nadeem et al., 2021; Parrish et al., 2022) and inherits the same vulnerability.

Blodgett et al. (2021) show that such benchmarks encode unexamined assumptions in their construction, so that part of what is measured is an artefact of how the data were built. Similarly, Se-shadri et al. (2022) show that altering the templates of such benchmarks in meaning-preserving ways changes the measured bias substantially.

Prior work validates the labels of constructed prompts, while we validate the construction itself, and find that the fillers needed to keep a template grammatical across all three stance versions are rarely stance-free (Section 2.2), and that this shifts stance estimates (Section 3.4.2).

LLM-generated evaluation data Generating synthetic data with LLMs has become a field of research on its own, raising many questions on the validity of such practice. Instruction-tuning pipelines generate prompts at scale from small seed sets (Wang et al., 2023; Ge et al., 2025) and work on simulated respondents generates whole populations of synthetic answers (Argyle et al., 2023), with later work questioning how faithfully these reproduce human distributions (Bisbee et al., 2024).

Prior work assesses whether synthetic (LLM-generated) data are valid against a label or a target distribution, not against what real users write. We generate prompts under detailed instructions, anchored in real prompts as seeds, and validate them against real prompts directly, asking humans and models which of three prompts a person plausibly typed, while also assessing how the intended factors are satisfied (Section 2.2).

Whether a model behaves differently once it recognises an evaluation has been studied mainly for capability benchmarks and generic transcripts. Van Der Weij et al. (2025) show that models can be made to underperform selectively. Similarly, Needham et al. (2025) find that frontier models can distinguish evaluation from deployment interactions, arguing that evaluation awareness is a source of distribution shift between testing and deployment.

Our realness ranking is a small instance of their classification task, applied to political prompts, and we find that templated prompts are separated from real ones more sharply, and more consistently, by our LLM annotators than by our human ones (Sec-