

tion 2.2). We do not establish that this recognition changes the responses themselves, and we say so in the Limitations.

LLM-as-a-judge for contested content Classifying the stance of a free-text model response requires a labeller, and LLM-as-a-Judge is a widely adopted paradigm (Zheng et al., 2023b; Gu et al., 2024). This choice holds its own limitations, with models known to favour their own outputs (Panickssery et al., 2024) and to be sensitive to the order in which candidates are presented. Röttger et al. (2026) use a single open-weight judge (Llama 3.1) carrying single-handedly the heavy task of identifying political leaning. On politically contested content, the risk is sharper than in a typical labelling task, since a model’s own leanings affect labelling.

We therefore use a majority-vote ensemble of three open-weight models developed under different regulatory environments, none of which shares a developer with the examined models, and we report their agreement per construction method and topic (Section 3.3).

6 Conclusion & Future Work

We set out to extend political stance detection beyond writing assistance, and found that the extension exposes a problem in the prompts themselves. Templated prompts, which work reasonably well for the writing-assistance requests they were designed around, degrade once the task becomes open-ended: their construction requires a filler that stays grammatical across all three stance versions, and such fillers are rarely stance-free. We proposed fully synthetic, LLM-generated prompts, anchored in real user prompts as seeds, as an alternative that preserves control over topic, intent, and user stance without inheriting that constraint.

In a validation study with three human and three LLM annotators, LLM-generated prompts are ranked on par with real prompts and above templated ones in how likely a real user is to type them, and carry their intended intent and stance more clearly. Their one weakness is topic under-specification in a few cases, which is repairable, whereas the defect of templated prompts is structural—and templating is also the construction our LLM annotators recognise most readily, which argues against templating given the evaluation-awareness concern we raised early on.

In a stance detection experiment with GPT-5.4 mini (and Grok 4.3 in Appendix G.1), the two

methods do not yield the same estimates; templated prompts constructed to convey a neutral user stance elicit sided responses on climate change and Russia–Ukraine, in the direction their fillers encode; under sided framings they shift the measured accommodation in whichever direction the filler points—understating it where filler and intended stance conflict and overstating it where they agree.

The outcome is that prompt construction is not an inconsequential design choice. Two studies of the same examined model, on the same topics, with the same judges, would report different political stances depending only on how their prompts were built, with the templated study overstating the model’s leanings most systematically where neutrality is what is being measured. We take this as an argument for treating prompt construction as a design choice in stance measurement rather than an implementation detail, and validating it explicitly.

In future work, we aim to further assess and improve the validity of the synthetic prompts with the ultimate goal of scaling up both the number of high-quality synthetic prompts and the number of examined topics, while also extending our work to other popular non-work-related intents, such as personal assistance, i.e., how the user should act in a given situation. We also aim to examine different phenomena, such as the temporal stance drift of a given model family, e.g., how GPT models change over time. We also aim to examine political stance under multi-turn dialogs to assess how follow-up interactions affect the model’s lean, i.e., whether the model becomes more or less accommodating of the user’s stance as the conversation develops.

Limitations

We report the core limitations of our work:

A narrow set of examined LLMs Our comparison of construction methods is run on two models from two US-based developers, OpenAI’s GPT 5.4 mini and xAI’s Grok 4.3 (Appendix G.1). We expect the mechanism (fillers that carry stance because they must stay grammatical across stance versions) to be model-independent, since it is a property of the prompts rather than of the model parsing them, and the two examined models do show the same qualitative pattern. The size of the gap we report, however, is not transferable, and a model with different alignment or refusal behaviour may absorb leading framings differently. Assessing how the difference carries across further models

and model families is left to future work.

LLMs are not AI chatbots The models we examine in our study are bare LLMs served via Application Programming Interfaces (APIs)¹⁹ without any additional wrapper or harness, including other modules, such as safety classifiers, geolocation (or other personalization) detectors, RAG support, access to the web or other tools, all of which affect how these systems respond as chatbots (GenAI assistants) when they sit behind commercial user interfaces. We expect the responses to differ, and leave a comparative analysis for future work.

A single generator model All synthetic prompts are generated by one flagship model (Claude Opus 4.8), which also performs the vagueness repair reported in Section 3. Our prompts therefore inherit that model’s stylistic tendencies, its notion of what a partisan user sounds like, and whatever coverage gaps its priors impose, none of which we isolate, although our generation instructions are designed to constrain them. Using a second generator would test how much of the “realness” gains are method-specific rather than model-specific. Moreover, blending collections of prompts generated by different models may be a mitigation mechanism to be considered in future work.

Scope of topics and language We cover 6 topics in English only, three of which are active geopolitical conflicts. Both the topic descriptions and the stance-loaded lexicons we supply reflect the debate as we understand it, and prompts and responses alike would look different if phrased differently. The chat logs we draw seeds and real prompts from (WildChat, LMSys-Chat) also over-represent the users of the platforms that produced them, at an earlier point in time. We do not assess how prompting in other languages, or framing the same issues within a specific national or regional context, e.g., the UK or the EU, affects the model’s responses. Our generation instructions ask for a global range of actors, places, and institutions (Appendix B), specifically instructing against defaulting to US/Western-centric prompts,²⁰ so our estimates speak to no particular country’s debate.

Detectability is not tested behaviourally We show that templated prompts are more readily iden-

tified as constructed, and more consistently so by the LLM annotators than by the human ones (Section 2.2), but not that this recognition changes how a model answers. Our annotators rank prompts explicitly; the examined models—the ones we collect responses with—were never asked to treat a prompt as an evaluation item, nor observed doing so. Establishing whether evaluation awareness alters stance would require probing the examined model itself, e.g., eliciting its own interpretation of a prompt’s source or inspecting its internal states (mechanisms), which we leave to future work.

LLM judges We validate the judges against each other rather than against human stance annotations of responses. The ensemble’s agreement therefore shows that the three judges converge, not that they converge on what a human annotator would label, i.e., how political stance is perceived by humans. Collecting human stance annotations for a subset of responses is left to future work. The broader problem, that a judge’s leanings are inseparable from its labels, is discussed in Section 4.

The average lean is a coarse measurement The lean averages an ordinal 5-point model judgment within a setting, so adjacent classes are treated as equally distant and a distribution of responses is reduced to its mean. Our statistical comparisons (Section 3.4.3) inherit this issue: they take the setting as the unit of analysis, over 18 neutral and 36 sided settings that are not fully independent, since they share the same six topics. A test over the individual judgments would be a better solution, but requires modelling the ordinal responses directly.

Acknowledgments

I would like to thank the National Center for AI in Society (CAISA) for providing the necessary funds to run the experiments presented in this study.

AI assistance

The author used Claude Code heavily for coding assistance on this project, while reviewing the code at all times. The author also used the Grammarly plugin and the Overleaf Writefull built-in assistant for spell-checking, grammar, and style editing, i.e., minor recommendation edits on sentence phrasing. In the final drafting of the paper, Claude Opus 5 was used for writing assistance and polishing.

¹⁹In our case, we use the OpenRouter API as a common access point for ease of development.

²⁰In early trials, we identified that Claude tended to generate mostly US-centric prompts.

References

- Steven Adler. 2025. [Five ways AI can tell you're testing it](#). Clear-Eyed AI (Substack).
- Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. 2023. [Out of one, many: Using language models to simulate human samples](#). *Political Analysis*, 31(3):337–351.
- Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. [On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?](#) In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, FAccT '21, page 610–623, New York, NY, USA. Association for Computing Machinery.
- Steven Bird. 2025. [Big AI is accelerating the metacrisis: What can we do?](#) *Preprint*, arXiv:2512.24863.
- James Bisbee, Joshua D. Clinton, Cassy Dorff, Brenton Kenkel, and Jennifer M. Larson. 2024. [Synthetic replacements for human survey data? the perils of large language models](#). *Political Analysis*, 32(4):401–416.
- Su Lin Blodgett, Gilsinia Lopez, Alexandra Olteanu, Robert Sim, and Hanna Wallach. 2021. [Stereotyping Norwegian salmon: An inventory of pitfalls in fairness benchmark datasets](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 1004–1015, Online. Association for Computational Linguistics.
- Maxwell T Boykoff and Jules M Boykoff. 2004. Balance as bias: Global warming and the us prestige press. *Global environmental change*, 14(2):125–136.
- Wayne Brittenden. 2026. [The Political Compass Test](#).
- Ilias Chalkidis and Stephanie Brandl. 2024. [Llama meets EU: Investigating the European political spectrum through the lens of LLMs](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*, pages 481–498, Mexico City, Mexico. Association for Computational Linguistics.
- Ilias Chalkidis and Anders Søgaard. 2026. [Brainrot: Deskillling and Addiction are Overlooked AI Risks](#), page 3005–3028. Association for Computing Machinery, New York, NY, USA.
- Aaron Chatterji, Thomas Cunningham, David J Deming, Zoe Hitzig, Christopher Ong, Carl Yan Shan, and Kevin Wadman. 2025. How people use ChatGPT? Technical report, National Bureau of Economic Research.
- Tavishi Choudhary. 2025. [Political Bias in Large Language Models: A Comparative Analysis of ChatGPT-4, Perplexity, Google Gemini, and Claude](#). *IEEE Access*, 13(0):11341–11379.
- Ricardo Dominguez-Olmedo, Moritz Hardt, and Celestine Mendler-Dünnér. 2024. [Questioning the survey responses of large language models](#). In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*.
- Esin Durmus, Liane Lovitt, Alex Tamkin, Stuart Ritchie, Jack Clark, and Deep Ganguli. 2024a. [Measuring the persuasiveness of language models](#).
- Esin Durmus, Karina Nguyen, Thomas Liao, Nicholas Schiefer, Amanda Askell, Anton Bakhtin, Carol Chen, Zac Hatfield-Dodds, Danny Hernandez, Nicholas Joseph, Liane Lovitt, Sam McCandlish, Orowa Sikder, Alex Tamkin, Janel Thamkul, Jared Kaplan, Jack Clark, and Deep Ganguli. 2024b. [Towards measuring the representation of subjective global opinions in language models](#). In *First Conference on Language Modeling*.
- Shangbin Feng, Chan Young Park, Yuhan Liu, and Yulia Tsvetkov. 2023. [From Pretraining Data to Language Models to Downstream Tasks: Tracking the Trails of Political Biases Leading to Unfair NLP Models](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11737–11762, Toronto, Canada. Association for Computational Linguistics.
- Jillian Fisher, Ruth Elisabeth Appel, Chan Young Park, Yujin Potter, Liwei Jiang, Taylor Sorensen, Shangbin Feng, Yulia Tsvetkov, Margaret Roberts, Jennifer Pan, Dawn Song, and Yejin Choi. 2025. [Position: Political neutrality in AI is impossible — but here is how to approximate it](#). In *Forty-second International Conference on Machine Learning Position Paper Track*.
- Tao Ge, Xin Chan, Xiaoyang Wang, Dian Yu, Haitao Mi, and Dong Yu. 2025. [Scaling synthetic data creation with 1,000,000,000 personas](#). *Preprint*, arXiv:2406.20094.
- Jiawei Gu, Xuhui Jiang, Zhichao Shi, Hexiang Tan, Xuehao Zhai, Chengjin Xu, Wei Li, Yinghan Shen, Shengjie Ma, and Honghao Liu. 2024. A survey on LLM-as-a-judge. *The Innovation*.
- William Guey, Pierrick Bougault, Vitor D. de Moura, Wei Zhang, and Jose O. Gomes. 2025. [Mapping Geopolitical Bias in 11 Large Language Models: A Bilingual, Dual-Framing Analysis of U.S.-China Tensions](#). *Preprint*, arXiv:2503.23688.
- Jochen Hartmann, Jasper Schwenzow, and Maximilian Witte. 2023. [The political ideology of conversational AI: Converging evidence on ChatGPT's pro-environmental, left-libertarian orientation](#). *Preprint*, arXiv:2301.01768.
- Noriaki Iwasa. 2010. The impossibility of political neutrality. *Croatian Journal of Philosophy*, 10(29):147–155.