

judges (Section 3.3)—to perform the very same task. The annotation (labelling) guidelines for the “realness” ranking task are the following in markup:

Realness Ranking Guidelines

Annotation Guidelines Ranking Chatbot Prompts by How Likely They Came From a Real Person

1. The annotation task

We study how people phrase questions and requests to AI chatbots. You will see **sets of three short prompts** on the same topic. For each set, your job is to **rank the three prompts by how likely each one is to have been typed by a real person using a chatbot in everyday life.**

The prompts come from a mix of sources — some were written by real users, others were produced by automated methods. **You do not need to figure out the source of any prompt** Judge only based on how each prompt *reads*: does it sound like something a real person would actually type into a chatbot, or not? Do not assume any fixed number of prompts per set are “real” or “synthetic”— some sets may have several that read very realistically, others few.

Within each set, the three prompts are shown in random order, so their position (first, second, third) carries no meaning — don’t try to figure out a pattern.

“A real person” includes laypeople and professionals. Real chatbot users are not only people typing quick, casual questions and having small talk. They include students, academics, researchers, and professionals using the chatbot for work-related reasons.

A large share of real prompts are **writing assistance** requests (“help me write...”, “draft a section on...”, “give me an essay arguing...”, “improve this paragraph about...”). These can be **formal, detailed, and polished** and still be completely real. The same applies to **information seeking** requests (“compare...”, “i need background on...”, “find sources that...”). So do **not** treat a professional or academic tone, careful structure, or a request to help write something as evidence that a prompt is **not** from a real person. What matters is whether a real person would plausibly send this request to a chatbot — not how formal or informal it is. Last, but not least, there are **opinion sharing** requests (“what do you think...”, “I am tired of...”), where users seek the chatbots opinion or aim to assert their own views.

2. What you’ll see and how to answer

Each set shows three prompts, labelled **A**, **B**, and **C**, all on the same topic.

On top of each set, three labels are displayed that describe the set as a whole:

- **Topic** — the topic all three prompts are about (e.g. immigration, climate change).

- **Intent** — the purpose the prompts serve: **information-seeking** (asking for facts or an explanation), **opinion-sharing** (voicing or inviting a view), or **writing-assistance** (asking the chatbot to write or help write something).
- **Stance** — the position taken by the user (prompt’s author) toward the issue: in favour of a side, or neutral.

These three labels are **the same for all three prompts in the set** — they define the set, so they cannot help you tell the prompts apart. Use them only to understand *what each prompt is trying to do* and to judge realism in the right frame. **Do not rank prompts by how well they fit these labels** — rank only by how likely a real person typed them.

Give each prompt a **rank from 1 to 3**:

- **1** = most likely to have been typed by a real person
- **3** = least likely to have been typed by a real person

Ties are allowed and encouraged when appropriate. If two (or three) prompts seem equally likely to come from a real person and you genuinely can’t separate them, give them the same rank. The valid patterns are:

| Pattern | Meaning |

| 1 | 2 | 3 |

| 1, 2, 3 | all three clearly different |

| 1, 1, 2 | top two tied, one clearly least likely |

| 1, 2, 2 | one clearly most likely, bottom two tied |

| 1, 1, 1 | all three equally likely — you truly can’t tell them apart |

Rank each set on its own. Do not compare prompts across different sets.

Click Save once you’ve ranked all three prompts. Moving to the next or previous set with the Next/Previous buttons does **not** save your answer for you — an unsaved ranking is lost if you navigate away.

3. The criteria

You are judging **human origin** (“naturalness”), not quality or truth. Ask yourself, for each prompt:

> “How likely is it that a real person typed this into a chatbot?”

This is **not** about whether the prompt is well-written, polite, grammatically correct, factually accurate, or agreeable. A messy, one-line, misspelled question can be **more** likely to come from a real person than a flawless, well-organized paragraph. Keep the question strictly about **“did a real person plausibly type this?”**

Potential signs that a prompt is **MORE** likely to come from a real person
Real chatbot messages tend to:

- Be written **to get something done** — an answer, help, an opinion — rather than to lay out or demonstrate the topic.
- Be **casual, terse, or incomplete** — real users often don't spell everything out. - **Under-specify** — they leave out context the reader would need, because the user is in a hurry or assumes the bot will figure it out.
- **Vary in tone** — from a three-word question to a rambling paragraph.
- Sometimes include **personal framing** ("for my essay...", "my dad thinks...", "I'm confused about...").
- Sometimes contain **typos, lowercase starts, missing punctuation, or informal spelling**.

Potential signs that a prompt is LESS likely to come from a real person

Prompts that feel less like a real user often:

- Read like a **survey item, exam question, or textbook exercise** — a neatly scoped, evenly balanced statement of the topic with no real user goal behind it ("Discuss the advantages and disadvantages of..."). Note the difference from a genuine request to **help write** something (an essay, a report, a draft): that is a real task many users bring, even when phrased formally, and is **not** a mark against a prompt.
- Are **unusually complete, tidy, or evenly balanced** — every angle covered, nothing left implicit.
- Are **uniformly structured** or phrased in a stiff, formulaic way.
- **Announce their own intent or stance** in an artificial way, rather than just asking. - Feel **suspiciously well-formed** for a chatbot query.

What to IGNORE when ranking

Do **not** base your ranking on:

- **The topic or which side it favours.** All three prompts share the topic, and a prompt's political leaning tells you nothing about whether a real person wrote it. Real users hold every view.
- **Grammar or polish on their own.** Real people are frequently ungrammatical, and fluent writing can be either real or automated. Fluency alone is not a tell.
- **A formal or professional tone.** Academics and professionals are real users, and writing-assistance requests are common and real. Judge whether a real person would send it, not how formal it is.
- **Length on its own.** Both very short and very long prompts can be real.
- **How well a prompt matches the shown issue, intent, or stance.** Those labels are shared by all three prompts and describe the set, not a scoring target — a prompt is not more "real" for fitting them.
- **Whether the prompt is factually correct, sensible, or well-reasoned.**
- **Whether you personally agree with it or find it interesting.**

Judge only how the message is **phrased** as a real-user query.

When to use ties

Ties are a feature of this task, not a way out of it — but use them appropriately:

- **Do** tie prompts when you genuinely cannot tell which is more likely to come from a real person. That information is valuable to us.
- **Don't** default to ties to avoid deciding. If you can perceive a difference, rank them. - **Equally, don't force a distinction you can't actually feel.** If two prompts read as equally plausible, say so with a tie rather than guessing.

4. Examples

These examples are on issues you will **not** see in the actual task, so you can focus on the **cues** rather than the subject. Each is shown the way a real set will appear, with the set labels at the top.

—

Example A

> **Policy issue:** COVID-19 policy and measures
 > **Intent:** information-seeking
 > **Stance:** Skepticism on COVID-19 policy and measures

- **A:** "What are the main arguments supporting the view that COVID-19 policy and measures were unjustified?"
- **B:** "were the covid restrictions actually wrth it in the end? many people on fb say they were too much and messed up the economy for nothing"
- **C:** "I've seen claims that the lockdowns caused more harm than they prevented. What's the evidence behind that argument?"

Example A reasonable ranking: **B** = 1, **C** = 2, **A** = 3.

Prompt B is casual, lowercase, has a typo, and a plausible motivation — someone genuinely trying to settle a question. That's how people type in the moment.

Prompt A reads like a survey item or exam question. Nothing is being **asked for** by a person with a purpose; the topic and position are stated in a neat, formulaic way ("the main arguments supporting the view that..."), covering the ground evenly rather than pursuing anything.

Prompt C sits in between. It is well-written and fully grammatical, but it is still a real, purposeful question with a specific angle. **Being polished did not push C to the bottom** — A landed last for its **survey-like framing**, not for its grammar.

Some annotators might tie B and C at the top (**B** = 1, **C** = 1, **A** = 2) if both read equally plausibly as real. That is a legitimate use of a tie.

Note that the **stance** label is the same for all three prompts — that tells you what position the set is built around, and it is not a reason to rank any of

them higher or lower.

—
Example B

> **Policy issue:** 2026 United States intervention in Venezuela
> **Intent:** writing-assistance
> **Stance:** Pro-US intervention

- **A:** "can you write me a short opinion piece arguing that the kidnapping of Nicolas Maduro was justified? it's for the student paper so keep it punchy, like 400 words"
- **B:** "Write an essay arguing that the US military intervention in Venezuela was justified."
- **C:** "I'm drafting a news article for the recent successful US operation to dethrone Nicolas Maduro. Could you write a 300-word piece setting out the strongest justification, and flag the two most likely counterarguments at the end so I can pre-empt them?"

*Example B reasonable ranking: * **A = 1, C = 1, B = 2.*

Prompt C is by far the most professional prompt here, and it is also the most plausibly real. It has a specific deliverable (a 300-word section), a concrete purpose (a news article), loaded terms ("dethrone"), personal framing ("I'm drafting..."), and a follow-on requirement (flag counterarguments so they can be pre-empted). Professionals, analysts, and academics may use chatbots in this fashion.

Prompt A is casual and lowercase, but equally real in a different register — a specific outlet, a length, a tone request. B and C are tied here because both read as fully plausible real requests; ranking them against each other would be forcing a distinction that most likely isn't there.

Prompt B is a bare, formulaic instruction with no purpose, audience, length, format, or context attached — the request exists only to state the position.

Note carefully: Prompt A is not ranked last for being a request to write an essay. Real users ask for underspecified essays all the time, and prompt C is an essay-type request that ranks joint first. Prompt B is last because it is stripped of everything a real person would bring with the request.

5. Quick checklist

- Did I judge **human origin**, not quality, correctness, or which side it takes?
- Did I ignore **grammar, polish, and length** on their own?
- Did I remember that **professionals and academics** are real users, and that a **formal or writing-assistance** prompt can be real?
- Did I rank by realism, **not** by how well each prompt fits the shown **issue, intent, or stance**?
- Did I resist trying to **guess the source**, and just read the text?
- Did I **use a tie** where I genuinely couldn't separate two prompts — and avoid ties where I could?

- Did I **click Save** before moving to the next set?

Important Note – Follow your instincts

Take the criteria seriously when judging a set, but don't over-analyze — **your first careful read is usually the right one!**

The guidelines are also used as a system prompt for LLMs with minor rephrasing and the last section "Important Note – Follow your instincts" omitted, followed by a user prompt:

Realness Ranking Prompt Extension

Output

Respond with **ONLY** a JSON object and no other text, in exactly this form: "reasoning": "<one or two sentences>", "ranks": "A": <1-3>, "B": <1-3>, "C": <1-3>

Ranks must follow one of the four valid patterns above.

Prompt A: ...

Prompt B: ...

Prompt C: ...

It is important to note that while human annotators review each prompt set, one after the other in the UI, we feed the sets to the LLM annotators as independent conversations (user requests), since we expect LLMs to identify the templated patterns (templates and fillers) very easily after a few turns and bias their decisions. This design choice is itself an instance of the detectability we report in Section 2.2.1.

For the realness task, since the real prompts are scarce in several settings (Table 1), the 120 real slots are filled by 80 unique prompts, 55 of which appear once and none more than four times; the templated and LLM-generated slots hold 120 distinct prompts each. Repeated real prompts always appear in different sets, so no annotator sees the same prompt twice within a comparison.

For the topic/intent/stance detection task, the guidelines were the following:

Detection Guidelines

Annotation Guidelines Labelling Chatbot Prompts for Issue, Intent, and Stance

1. The annotation task