

two human annotators is almost perfect on topic ($\kappa=.96$), while being substantial on intent ($\kappa=.63$) and stance ($\kappa=.76$) following Landis and Koch (1977). As we describe in our findings, annotators were heavily challenged by value-laden questions, which they did not consistently treat as calling for a value judgment (opinion); this accounts for the lower agreement on intent. The LLM annotators align almost perfectly across all variables: topic ($\kappa=.96$), intent ($\kappa=.83$), and stance ($\kappa=.86$).

G Additional Results

G.1 Stance Analysis on Grok 4.3

Following our analysis on stance detection with OpenAI GPT 5.4 mini responses for both prompt construction methods, in Table 10 we present results for xAI’s Grok 4.3 using the same prompts and the same judge ensemble.

General effects Considering general observations of the model’s stance, irrespective of the set of prompts, we observe that Grok follows a similar pattern, leaning towards a climate urgency position, but to a slightly lesser extent in some cases, e.g., 0.45 difference in neutral opinion sharing, compared to GPT 5.4 mini. Grok 4.3 can better follow a climate-moderation stance in writing assistance requests (1.23-1.56), compared to GPT 5.4 mini (0.73-1.00), but is equally “unwilling” to accommodate a climate-moderate stance in opinion-sharing. On immigration, similar to GPT 5.4 mini, Grok 4.3 stays close to neutral, and accommodates an immigration-expansive stance more than its counter on information seeking (-1.35/-0.52 against 0.51/0.07), but not on opinion sharing, where it accommodates neither pole (-0.24/-0.20 against -0.14/0.00), unlike GPT 5.4 mini. On AI adoption, the two models are closest: both stay at neutral under neutral framings, and both accommodate the user in either direction, so AI adoption remains the one topic on which accommodation runs both ways.

On Israel–Palestine, Grok 4.3 accommodates a pro-Israel stance considerably more on information seeking and writing assistance requests, and is less “willing” to support a pro-Palestine stance on opinion sharing (0.15/0.06 against 0.53/0.65 for GPT-5.4 mini). Under neutral framings, neither model holds a pro-Palestine stance. On Russia–Ukraine, Grok 4.3 has a very similar picture to GPT-5.4 mini; substantially leaning towards a pro-Ukraine stance. On US/Israel–Iran, both models

stay neutral under neutral framings and resist a pro-US/Israel framing on the two non-writing intents (-0.16/-0.20 for Grok 4.3), while moving toward pro-Iran positions to a lesser extent than GPT-5.4 mini (0.72/0.21 against 1.16/0.23 in information seeking). In writing assistance, however, Grok 4.3 complies substantially more with pro-US/Israel requests (-1.19/-1.77 against -0.41/-1.10).

Comparative Cross-Method Analysis Considering the differences between the two prompt construction methods, we observe a very similar pattern. The templated prompts tend to lead to more polarized responses. In the topic of climate change, the neutral templated prompts lead to a substantially larger lean towards the climate urgency position (-1.25/-0.27/-1.36 differences across the three examined intents). In the topic of Israel–Palestine, the gap between templated and LLM-generated prompts closes, given Grok’s general stance. In the topic of Russia–Ukraine, similar to GPT-5.4 mini, Grok is less willing to accommodate a pro-Russia stance when using templated prompts compared to LLM-generated ones (1.22/1.42 against 0.12/0.35 in information seeking and opinion sharing). The two topics we did not examine above show the same filler-induced siding we report in Section 3.4.2: on AI adoption, the templated fillers tilt toward AI-scepticism in information seeking, weakening the accommodation of AI-positivism (-0.38 against -1.23) and strengthening that of AI-scepticism (1.48 against 0.39), a wider spread than the one we observe for GPT-5.4 mini; on US/Israel–Iran, the pro-Iran filler compounds the siding that “the US and Israeli strikes on Iran” already carries, and templated information-seeking prompts elicit a clear pro-Iran lean (0.72) where the LLM-generated ones stay close to neutral (0.21).

The replication is close in direction, if not always in magnitude: across the 18 neutral settings, responses to templated prompts sit 0.36 from neutral against 0.05 for LLM-generated ones (Wilcoxon signed-rank, $p < .001$; templated further from neutral in 15 of the 18 settings), compared with 0.48 and 0.07 for GPT-5.4 mini. Under sided framings, the two methods diverge by 0.46 scale points per setting on average, against 0.41 for GPT-5.4 mini, and again not in a consistent direction.

Summary First, the effect of the construction method is not specific to a single model or developer: on both examined models, templated prompts are systematically further from neutral in the di-

rection their fillers encode, and under sided framings they shift the measured accommodation in whichever direction the filler points. Second, the models’ own stances are not the same, and our results should not be read as if they were: Grok 4.3 accommodates pro-Israel framings more and pro-Palestine ones less, complies more readily with pro-US/Israel and climate-moderate writing requests, and accommodates immigration-expansive framings less in opinion sharing. The observation we draw in Section 3.4.1, that the model holds a consistent stance against the actor who initiates a military operation, therefore holds for Grok 4.3 on Russia–Ukraine but is weaker on Israel–Palestine. What is consistent across the two models is the effect of prompt construction, not the leanings it measures.

G.2 LLM Judge Models

As we reported in Section 3.3, the agreement among the three judges is high for both construction methods, with an ordinal Krippendorff’s α of .91 and .82 for responses to templated and LLM-generated prompts, respectively. DeepSeek V4 Pro is the judge model that deviates (disagrees) the most on its stance judgments, compared to the other two, i.e., Mistral Large 3 and NVIDIA’s Nemotron 3 Ultra agree with each other more than either does with DeepSeek. The deviation is slight in the pairwise scores: .90/.80 (vs Mistral) and .90/.84 (vs. Nemotron) in ordinal Krippendorff’s α for templated and LLM-generated, with the other two scoring .92/.84 (Mistral vs. Nemotron).

Considering agreement per topic, we find that the lowest agreement is on the topic of Russia–Ukraine: .60 and .72 agreement on templated and LLM-generated prompts, respectively. Similarly, for climate change, the agreement is .71 and .76, respectively. When we also consider individual intents per topic, we find that opinion-sharing templated prompts have the lowest agreement at .53. Nonetheless, the effect of potential disagreements is not visible in our analysis (Section 3), since we consider a three-way majority-vote ensemble, excluding all samples, where such a majority cannot be established (43 and 69 responses, approx. 2% and 3% of all responses, for templated and LLM-generated prompts, respectively). LLM judge models’ disagreement and its causes are a topic that we leave to future work.

In Table 11, we present the difference between the DeepSeek model and the majority-vote ensemble on OpenAI GPT 5.4 mini responses. Over-

all, the differences are small, with a mean absolute difference of 0.07 across all settings and a maximum of 0.32. In general, DeepSeek shifts slightly towards an immigration-expansive (-0.03), AI-positivism (-0.04), pro-Palestine (+0.03), pro-Russia (-0.04), and pro-US/Israel position (-0.04), and no shift on climate change (+0.00).

Issue	Pole A	Pole B
AI adoption and policy	<i>AI-positivism</i> – innovation and diffusion. Benefits (productivity, scientific discovery, accessibility to expertise) are large and near-term; harms are speculative, manageable, or outweighed. Favours permissive deployment and voluntary standards.	<i>AI-scepticism</i> – precaution and accountability. Harms (labour displacement, surveillance, discrimination, misinformation, concentration of corporate power, safety failures) are already material. Favours binding obligations on developers and providers.
Climate change and action	<i>Climate-urgency</i> – rapid, state-led decarbonisation. Warming is severe, human-caused, and near a threshold; treats climate as an overriding policy priority. Favours binding targets, public investment in renewables, and fossil-fuel phase-out.	<i>Climate-moderation</i> – growth, adaptation, and cost realism. Emphasises the economic and distributional costs of transition, energy security and affordability, the limits of unilateral action, and doubts about projected severity. Favours technology-neutral incentives and adaptation.
Immigration	<i>Immigration-expansive</i> – openness and migrant rights. Immigration is net-positive economically and culturally enriching; humanitarian duties to migrants and refugees can override border control, and freedom of movement is treated as a right. Favours accessible legal pathways and strong protections.	<i>Immigration-restrictive</i> – control and cohesion. Border control is a core sovereign function; emphasises pressure on public services, housing, and wages, rule-of-law concerns around irregular entry, cultural continuity, and demographic change as unmanaged or deliberate. Favours lower intake and stricter enforcement.
Israel–Palestine	<i>Pro-Israel</i> – security and self-defence. Foregrounds Hamas’s 7 October attack and hostage-taking, and Israel’s right to defend itself and eliminate the military threat; attributes civilian harm to Hamas’s use of civilian areas. Favours continued operations until security objectives are met.	<i>Pro-Palestine</i> – self-determination and rights. Foregrounds occupation, settlement expansion, blockade, and displacement, and the scale of Palestinian civilian death, starvation, and destruction in Gaza; treats Israeli conduct as disproportionate. Favours cease-fire, statehood, and accountability.
Russia–Ukraine	<i>Pro-Russia</i> – security interests and negotiated settlement. Foregrounds NATO enlargement, Western involvement in Ukrainian politics, and Russian security claims as causes; questions the cost and escalation risk of Western support. Favours a negotiated end, including territorial concessions.	<i>Pro-Ukraine</i> – sovereignty and Western support. An unprovoked war of aggression against a sovereign state; territorial integrity and the norm against conquest are at stake. Favours sustained military and financial aid, sanctions, and accountability.
US/Israel–Iran	<i>Pro-US/Israel</i> – non-proliferation and deterrence. Foregrounds Iran’s enrichment beyond civilian need and IAEA non-compliance, its proxies, and its threats against Israel; treats strikes as preventive. Favours sanctions and military pressure, including support for regime change.	<i>Pro-Iran</i> – sovereignty and non-intervention. Foregrounds strikes on Iranian territory as violations of international law, the legality of a civilian programme under the NPT, Western double standards, and the human cost of sanctions. Favours diplomacy and recognition of Iran’s regional role.

Table 5: Description of the examined topics, alongside the two opposing poles.